

CS3.401 Distributed Systems

Name- Vilal Ali

Roll Number: 2024701027

2024701027_Homework_1

Question 1. Consider a system with four processors and a distributed program with about 15 events.

Question 1 (a) - Draw the time diagram for such a program and mark the scalar time of the events. The set of events should include all possible events.

Answer:

Let's consider a distributed system of four processes (P1, P2, P3, P4) with 15 events. Scalar timestamps are assigned to each event using Lamport's logical clocks, where each event within a process receives a timestamp that increases in a sequential manner.




Figure 1.0 The space-time diagram of a distributed execution.

Scalar Time:

Below is the time space diagram created to represent this. Here we are set the increment d to 1 always.




Figure 1.1: The space-time diagram of a distributed execution.

Question 1 (b) - Identify events that are logically concurrent:

Answer:

- e_i^x : Represents an event x on process P_i .
- e_j^y : Represents an event y on process P_j .

Two events e_i^x and e_j^y are logically concurrent if there is no path (causal relationship) between them.

Logically Concurrent Events Are:

- $e_1^1 || e_2^1, e_1^1 || e_3^1, e_1^1 || e_4^1$
- $e_2^1 || e_3^1, e_2^1 || e_4^1$
- $e_3^1 || e_4^1, e_2^2 || e_4^2, e_3^2 || e_4^2$

Question 1 (c) - Identify events where strong monotonicity of the scalar time fails.

Answer: Scalar time should preserve the "happens-before" relation (\rightarrow), meaning that if $e_i \rightarrow e_j$ then $C(e_i) < C(e_j)$.

In the diagram, strong monotonicity of scalar time could fail if there is any sequence where an event e^x_i occurs after another event e^y_j , but $C(e^x_i)$ is less than $C(e^y_j)$. Without actual clock values provided, we can analyze the sequence of events visually:

From Figure (1.1), all the direct causality links (arrows) maintain the order of events without causing any strong monotonicity failures.

Question 2 (a): Draw the time diagram for such a program and mark the vector time of the events. The set of events should include all possible events.

Answer: Vector time ensures that the "happens-before" relation is preserved across all processes with vector clocks:

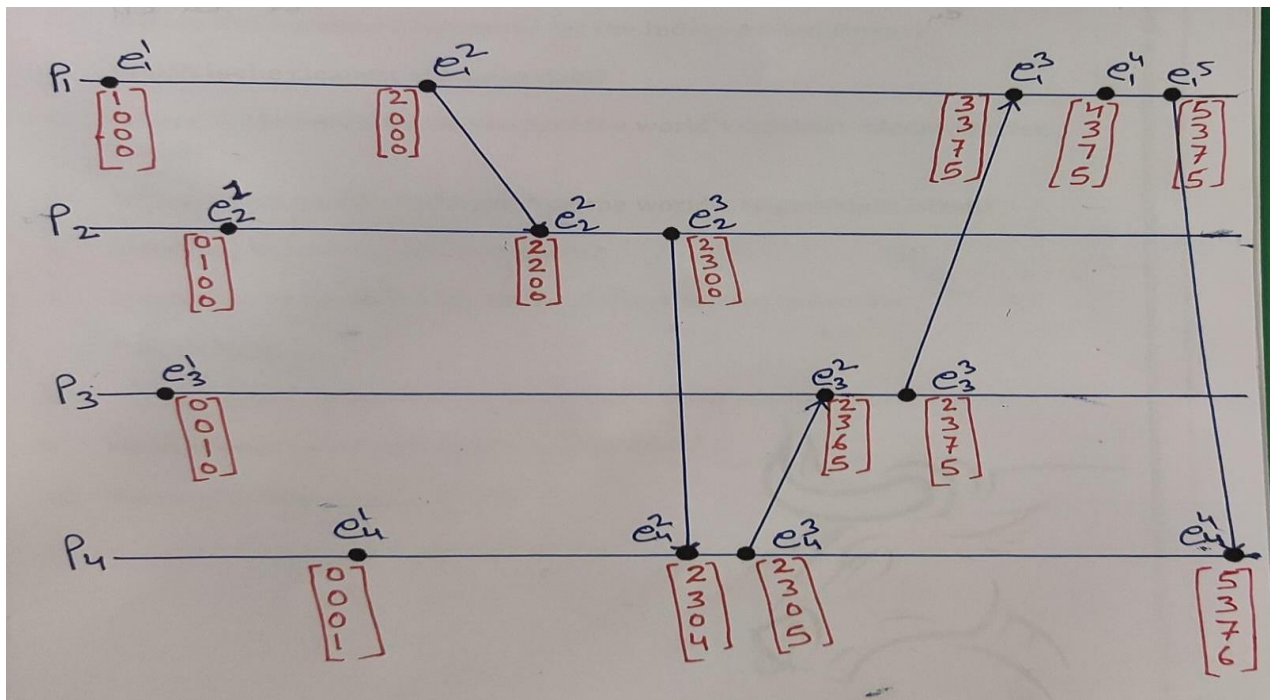


Figure 2.0: The vector space-time diagram of a distributed execution.

- Initially, the vector clock is set to $[0, 0, 0, \dots, 0]$ and $d = 1$.
- Assign timestamps to all the events shown in the diagram:

For each event, we have the following vector timestamps:

P1	P2	P3	P4
e^1_1 P1: [1,0,0,0]	e^1_2 P2: [0,1,0,0]	e^1_3 P3: [0,0,1,0]	e^1_4 P4: [0,0,0,1]
e^2_1 P1: [2,0,0,0]	e^2_2 P2: [2,2,0,0]	e^2_3 P3: [2,3,6,5]	e^2_4 P4: [2,3,0,4]
e^3_1 P1: [3,3,7,5]	e^3_2 P2: [2,3,6,5]	e^3_3 P3: [3,3,7,5]	e^3_4 P4: [2,3,0,5]
e^4_1 P1: [4,3,7,5]			e^4_4 P4: [5,3,7,6]

e^5_1 P ₁ : [5,3,7,5]			
------------------------------------	--	--	--

Table: 1.0: Given Events and Vector Timestamps

Question 2 (b): Identify events that are logically concurrent with respect to vector time.

Answer:

To identify logically concurrent events in the second image, we need to determine pairs of events that have no causal relationship. This means there is no direct or indirect "happens-before" relationship between them.

In the context of vector clocks, two events are considered logically concurrent if their vector timestamps are not comparable, meaning that neither event's timestamp is a component-wise greater than or equal to the other event's timestamp.

Identifying Logically Concurrent Events:

Process P₁:

- e^1_1 is logically concurrent with:
 - e^1_2 [0, 1, 0, 0] (P₂)
 - e^1_3 [0, 0, 1, 0] (P₃)
 - e^1_4 [0, 0, 0, 1] (P₄)
- e^2_1 is logically concurrent with:
 - e^1_2 [0, 1, 0, 0] (P₂)
 - e^1_3 [0, 0, 1, 0] (P₃)
 - e^1_4 [0, 0, 0, 1] (P₄)
- e^3_1 is logically concurrent with:
 - e^2_2 [2, 2, 0, 0] (P₂)
 - e^2_4 [2, 3, 0, 4] (P₄)

Process P₂:

- e^1_2 is logically concurrent with:
 - e^1_1 [1, 0, 0, 0] (P_1)
 - e^1_3 [0, 0, 1, 0] (P_3)
 - e^1_4 [0, 0, 0, 1] (P_4)
- e^2_2 is logically concurrent with:
 - e^3_1 [3, 0, 0, 0] (P_1)
 - e^2_3 [2, 3, 5, 0] (P_3)
 - e^2_4 [2, 3, 0, 4] (P_4)

Process P_3 :

- e^1_3 is logically concurrent with:
 - e^1_1 [1, 0, 0, 0] (P_1)
 - e^1_2 [0, 1, 0, 0] (P_2)
 - e^1_4 [0, 0, 0, 1] (P_4)
- e^2_3 is logically concurrent with:
 - e^2_4 [2, 3, 0, 4] (P_4)
- e^3_3 is logically concurrent with:
 - e^4_1 [3, 3, 7, 5] (P_1)
 - e^3_4 [2, 3, 0, 5] (P_4)

Process P_4 :

- e^1_4 is logically concurrent with:
 - e^1_1 [1, 0, 0, 0] (P_1)
 - e^1_2 [0, 1, 0, 0] (P_2)
 - e^1_3 [0, 0, 1, 0] (P_3)
- e^2_4 is logically concurrent with:
 - e^3_1 [3, 0, 0, 0] (P_1)
 - e^2_2 [2, 2, 0, 0] (P_2)
 - e^2_3 [2, 3, 5, 0] (P_3)

Summary of Logically Concurrent Events:

- e^1_1 with e^1_2, e^1_3, e^1_4

- e^2_1 with e^1_2, e^1_3, e^1_4
- e^3_1 with e^2_2, e^2_4
- e^1_2 with e^1_1, e^1_3, e^1_4
- e^2_2 with e^3_1, e^2_3, e^2_4
- e^1_3 with e^1_1, e^1_2, e^1_4
- e^2_3 with e^2_4
- e^3_3 with e^4_1, e^3_4
- e^1_4 with e^1_1, e^1_2, e^1_3
- e^2_4 with e^3_1, e^2_2, e^2_3

These events are logically concurrent because they have no causal dependency on each other, as indicated by their vector timestamps.

Question 3: Show that the timestamps assigned by vector time are strongly consistent.

Answer: To demonstrate that the timestamps assigned by vector time are strongly consistent, we need to show that for any two events e_i and e_j in a distributed system:

If $e_i \rightarrow e_j$ (i.e., event e_i happens-before event e_j), then the vector timestamp of e_i is less than or equal to the vector timestamp of e_j in every component.

Note: This is the condition of strong consistency in the context of vector clocks.

Analysis of Event Causality: Using the Table 1.0

1. Causality within a Process:

- Within any single process P_i , events are ordered in time by their index. For example:
 - $e^1_1 \rightarrow e^2_1 \rightarrow e^3_1 \rightarrow e^4_1 \rightarrow e^5_1$
- This implies that the vector timestamps must be monotonically non-decreasing:
 - $[1, 0, 0, 0] \leq [2, 0, 0, 0] \leq [3, 0, 0, 0] \leq [3, 3, 7, 5] \leq [5, 3, 7, 5]$
- We see that this holds true for each process.

2. Causality across Processes:

- Consider the causal relationships between events across different processes, as indicated by the arrows in the diagram. For example:
 - $e^2_1 \rightarrow e^2_2$
 - Vector timestamps: $[2, 0, 0, 0] \leq [2, 2, 0, 0]$
 - $e^2_2 \rightarrow e^2_3$
 - Vector timestamps: $[2, 2, 0, 0] \leq [2, 3, 5, 0]$
 - $e^3_2 \rightarrow e^3_3$
 - Vector timestamps: $[2, 3, 6, 5] \leq [2, 3, 7, 5]$
 - $e^2_4 \rightarrow e^3_4$
 - Vector timestamps: $[2, 3, 0, 4] \leq [2, 3, 0, 5]$

These relations hold true according to the vector timestamps. This guarantees that vector time is strongly consistent, meaning it correctly captures the causal relationships between events in a distributed system.