

---

Distributed Systems

Monsoon 2021

Lecture 24

International Institute of Information Technology

Hyderabad, India

---

# Quick Recap

---

- Started studying distributed algorithms, mostly on graphs.
- Looked at problems such as
  - BFS
  - Spanning Tree (Asynchronous)
  - Shortest Paths.
  - Primitives such as Convergecast and Broadcast
  - MST

# Maximal Independent Sets (MIS)

---

- For a graph  $(V, E)$ , an independent set of nodes  $U$ , where  $U \subseteq V$ , is such that for each  $i$  and  $j$  in  $U$ ,  $(i, j) \notin E$ .
- An independent set  $U$  is a **maximal independent set** if no proper superset of  $U$  is an independent set.
- A graph may have multiple MIS; perhaps of varying sizes.
- The largest sized independent set is the maximum independent set.
  - Application: wireless broadcast - allocation of frequency bands (mutex)
  - NP-complete

# Maximal Independent Sets (MIS)

---

- For a graph  $(V, E)$ , an independent set of nodes  $U$ , where  $U \subseteq V$ , is such that for each  $i$  and  $j$  in  $U$ ,  $(i, j) \notin E$ .
- An independent set  $U$  is a **maximal independent set** if no proper superset of  $U$  is an independent set.
- Has been a very popular problem with ongoing research in distributed algorithms.

# Maximal Independent Sets (MIS)

---

- Has been a very popular problem a problem with ongoing research in distributed algorithms.
- A greedy sequential algorithm is very easy to design.

# Greedy Sequential Algorithm

---

Algorithm GreedyMIS(G)

Begin

$U = \{\}, S = V$

While (S is not empty) do

    Pick a vertex  $v$  in  $S$  and add  $v$  to  $U$

    Remove  $v$  and all neighbors of  $v$  in  $S$

End-While

Return  $U$

End.

# Maximal Independent Sets (MIS)

---

- Has been a very popular problem a problem with ongoing research in distributed algorithms.
- A greedy sequential algorithm is very easy to design.
- But such an algorithm can take a long time in the distributed setting.
- Need quick way to set apart neighbors so that no two neighbors join the MIS together.
- Problem is called as **symmetry breaking**.

# Maximal Independent Sets (MIS)

---

- Need quick way to set apart neighbors so that no two neighbors join the MIS together.
- Problem is called as symmetry breaking.
- In the deterministic setting, it is shown that any algorithm will require nearly  $n^\epsilon$  rounds.
- Randomization however is very helpful.



# Luby's Algorithm for MIS

---

- The algorithm designed by Luby more than three decades ago works as follows.
- Every node maintains a state among `{active, inactive}`.
- Every active node  $v$  picks a random number 0 or 1 with probability  $\frac{1}{2} d(v)$ .
  - 1 means it is interested in joining the MIS.
- Each such node then shares its choice with all its active neighbors.
- A node withdraws from the contest if it finds another neighbor of **high degree** that also chose 1.
- Such a withdrawing node still remains `active` for now.

# Luby's Algorithm for MIS

---

- Every node maintains a state among `{active, inactive}`.
- Every active node picks a random number 0 or 1.
- Each such node then shares its choice with all its active neighbors.
- A node withdraws from the contest if it finds another neighbor of high degree that also chose 1. Such a withdrawing node still remains `active` for now.
- If an active node that chose 1 did not withdraw, then the node joins the MIS. Becomes `inactive`.
  - Indicates all its active neighbors of its decision to join MIS.
- Active nodes that have some neighbor joining an MIS turn `inactive`.

# Luby's Algorithm for MIS

---

- It is not immediately clear as to how this randomization helps.
- It can be shown with some analysis that the algorithm requires  $O(\log n)$  rounds with high probability.
- The above is shown by arguing that in each round, a fraction of the edges can be made redundant.
  - An edge is redundant if either of its end points are `inactive`.
  - This indicates that in  $O(\log n)$  rounds, all edges are redundant.
- The analysis follows....

# Analysis

---

- To carry out the analysis, define:
  - A vertex  $v$  is **good** if more than  $1/3^{\text{rd}}$  of its neighbors have degree less than  $d(v)$ .
  - A vertex  $v$  is **bad** if at least  $2/3^{\text{rd}}$  of its neighbors have degree at least  $d(v)$ .
  - An edge  $e$  is **good** if at least one endpoint of  $e$  is **good**.
  - An edge  $e$  is **bad** if both its endpoints are **bad**.

# Analysis – Sketch

---

- Good vertices have enough low degree neighbors.
- The probability that a node chooses to join the MIS is inversely proportional to its degree.
  - Low degree is good!
- So at least one such low degree neighbor is in  $S$  with good probability.
  - Helps delete good vertices.
  - This in turn helps delete good edges.
- If we can show that there are enough good edges, then suffices if a fraction of them are deleted.

## Analysis – Claim 1

---

- For every good vertex  $v$  with  $d(v) > 0$ , the probability that some neighbor  $w$  of  $v$  gets marked is at least  $1 - e^{-1/6}$
- Proof:  $\Pr(v \text{ is marked}) = \frac{1}{2}d(v)$ .
- Since  $v$  is good, at least  $d(v)/3$  neighbors have degree at most  $d(v)$ . Let  $w$  be such a neighbor.
- $\Pr(w \text{ is marked}) = \frac{1}{2} d(w) \geq \frac{1}{2}d(v)$ .
- $\Pr(\text{no low degree neighbor of } v \text{ is marked}) \leq \left(1 - \frac{1}{2}d(v)\right)^{d(v)/3} = e^{-1/6}$ .
- We did not consider the actions of the high degree neighbors in the above calculation.

## Analysis – Claim 2

---

- If a vertex  $w$  is marked, then  $\Pr(w \text{ in } S) \geq \frac{1}{2}$ .
- Proof: If  $w$  is marked, then  $w$  is not in  $S$  only if some high degree neighbor of  $w$  is also marked.
- Each such high degree neighbors of  $w$  is marked with probability at most  $\frac{1}{2d(w)}$ .
- Number of such high-degree neighbors  $\leq d(w)$ .

$$\begin{aligned}\Pr(w \text{ in } S \mid w \text{ is marked}) &= 1 - \Pr(w \text{ not in } S \mid w \text{ is marked}) \\ &= 1 - \Pr(\exists u \in N(w), d(u) \geq d(w), u \text{ marked}) \\ &\geq 1 - |u \in N(w), d(u) \geq d(w)| \cdot \frac{1}{2d(w)} \\ &\geq 1 - |u \in N(w)| \cdot \frac{1}{2d(w)} = 1 - \frac{d(w)}{2d(w)} \\ &= \frac{1}{2}.\end{aligned}$$

## Analysis – Claim 3

---

- Let  $v$  be a good vertex. Then,  
 $\Pr(v \text{ is deleted}) \geq (1 - e^{-1/6})/2.$
- Proof: Combine Claims 1 and 2.




## Analysis – Claim 4

---

- At least half the edges are good.
- Proof: For every bad edge  $e$ , associate a pair of edges via a function  $f:E_B \rightarrow \binom{E}{2}$  such that for any two distinct bad edges  $e_1$  and  $e_2$ ,  $f(e_1) \cap f(e_2) = \emptyset$ .
- Completes the proof since only  $|E|/2$  such pairs exist.
- The function  $f$  defined as follows.

# Analysis – Claim 4

---



- For each edge  $(u, v) \in E$ , orient it towards the vertex of higher degree.
- Consider a bad edge  $e = (u, v)$  oriented towards  $v$ .
- Since  $v$  is bad, the out-degree of  $v$  is at least twice its in-degree.
- So, there exists a way to pick a pair of edges for every bad edge.

# Putting Everything Together

---

- In each iteration, it is expected that a constant fraction of edges are deleted.
  - Half the edges are good, and a good edge is deleted with probability at least  $(1 - e^{-1/6})/2$ .
- So, on expectation, only  $O(\log m) = O(\log n)$  iterations suffice.
- Can also show that with high probability  $O(\log n)$  iterations suffice.


# Leader Election

---

- Another problem that is studied in the distributed setting is leader election.
- The problem is defined as follows.
- Let  $G$  be a (directed) graph  $(V, E)$ . One of the nodes of  $V$  has to be designated as the leader.
- Several applications
  - We saw an application to the 2-phase commit algorithm also.

# Leader Election


---



- As a simple example, consider the case where  $G$  is a (bidirectional) directed ring.
- Processes can distinguish clockwise from counterclockwise neighbor.

# Leader Election


---



- General idea:
  - Every process passes its UID around the ring
  - Upon reception, compare to own UID
    - If larger: Pass to next process
    - If smaller: Discard
    - If equal: Declare self as leader

# Leader Election


---



- Proof by induction on  $r$ 
  - After  $r$  rounds, the node  $r$  away from the  $P_{max}$  will be sending  $UID_{max}$
  - When  $r = n$ ,  $UID_{max}$  will have traveled around the ring

# Leader Election

---




- An improvement is given by Hirshberg and Sinclair
- Achieves  $O(n \log n)$  communication complexity,  $O(n)$  time complexity



# Leader Election


---



- General idea:
  - Operates in phases
  - During phase  $p$ , send UID in both directions for a distance of  $2^p$ 
    - Processing of received UID same as earlier
    - If its UID comes back, continue to next round
    - If it receives its outbound UID, declare self as leader

# General Graphs


---



- For general graphs, one of the popular algorithms is the FloodMax algorithm.
- General idea:
  - Each process remembers the largest UID it has seen
  - For each round, propagate  $UID_{\max}$  on all outgoing links
  - After *diam* rounds, if  $UID_{\max}$  is own UID, declare self as leader

# General Graphs

---



- Communication complexity is  $diam \cdot |E|$ 
  - $|E|$  is the number of outgoing edges