

# The Linear MPC Model

---

- Having  $S = n^{1+\varepsilon}$  makes it very easy to design algorithms that are super-fast.
- Let us reduce  $S$  to  $\Theta(n)$ . We call this the linear MPC model.
- In the linear model, each machine has enough space to store a very limited information about each node in the graph.
- Let us continue the MST example and see how to design algorithms in this model.

# The Linear MPC Model

---

- The techniques needed here are quite involved.
- Consider a sorted order of edges by weight. Let the ordering be  $e_1, e_2, \dots, e_m$ .
- The following observation holds:

Edge  $e_i$  is in the MST if and only if its endpoints are not in the same connected components in the graph with edge set  $\{e_1, \dots, e_{i-1}\}$ .
- We can use this observation directly to convert the MST problem to the connected components problem.
  - Check the observation for each of  $e_1$  to  $e_m$ .
- However, this algorithm is quite wasteful.

# The Linear MPC Model

---

- One improvement is as follows.
- Group the  $m$  edges into  $m/n$  groups of  $n$  edges each.
- Call  $E_i = \{e_{(i-1)n+1}, \dots, e_{in}\}$  for  $i = 1$  to  $m/n$ .
- Let  $H_i = \bigcup_{j=1}^i E_j$ .  $H_i$  is the union of the first  $i$   $E_i$ s for  $i=1$  to  $m/n$ .
- Let  $F_i$  be the MST (Forest) corresponding to  $H_i$ .
- Notice that  $|F_i| = O(n)$ .
- So, given  $F_i$  and  $E_{i+1}$ , both of size  $O(n)$  edges, one machine can check which edges in  $E_{i+1}$  will be in the MST.
- We could compute each  $F_i$  in parallel!

# The Linear MPC Model

---

- Call  $E_i = \{e_{(i-1)n+1}, \dots, e_{in}\}$  for  $i = 1$  to  $m/n$ .
- Let  $H_i = \bigcup_{j=1}^i E_j$ .  $H_i$  is the union of the first  $i$   $E_j$ s for  $i=1$  to  $m/n$ .
- Let  $F_i$  be the MST (Forest) corresponding to  $H_i$ .
- Notice that  $|F_i| = O(n)$ .
- So, given  $F_i$  and  $E_{i+1}$ , both of size  $O(n)$  edges, one machine can check which edges in  $E_{i+1}$  will be in the MST.
- We could compute each  $F_i$  in parallel!
- But,  $|H_i|$  is too big to fit in any single machine!!

# The Linear MPC Model

---

- We just wish that computing the  $F_i$ s takes as few rounds as possible.
- Here is a way to do that.
- We use the technique of graph sketching.
- We then explain how to implement the technique in the linear MPC model to compute connected components in  $O(1)$  rounds.

# The Linear MPC Model

---

- Let  $A$  be a subset of  $V$ . Let  $C$  be the cut between  $A$  and  $V \setminus A$ .
- Assume for a moment that the cut has only edge crossing it.
- We wish to find the cut edge across  $A$  and  $V \setminus A$  using only  $O(\log n)$  bits of memory.
- We use a bit trick that works as follows.


# The Linear MPC Model

---

- Identify each edge by the concatenation the identifiers of its endpoints. For  $e = uv$ ,  $\text{id}(e) = u \circ v$  if  $\text{id}(u) < \text{id}(v)$ .
- Locally, each vertex  $u$  computes 
$$\text{XOR}_u = \bigoplus_{e \in E, e \text{ has } u \text{ as an end point}} \text{id}(e).$$
- Now, consider the XOR of all the vertices in  $A$ , 
$$\text{XOR}_A = \bigoplus_{u \in A} \text{XOR}_u.$$
- In  $\text{XOR}_A$ , an edge  $e = uv$  with both end points  $A$  appears in both  $\text{XOR}_u$  and  $\text{XOR}_v$ .
- On the other hand, the edge  $e = uv$  with  $u$  in  $A$  and  $v$  in  $V \setminus A$ , appears only once (in  $\text{XOR}_u$ ).
- So, the result of  $\text{XOR}_A = \text{id}(e)$  thereby allowing us to identify the cut edge.

# The Linear MPC Model

- Example



Edge	End points	ID	Node	Edges	XOR <sub>v</sub>
e1	a, c	000010	a	e1, e2	000011
e2	a, b	000001	b	e2, e3	001010
e3	b, d	001011	c	e1, e4, e8	000101
e4	c,d	010011	d	e3, e4	011000
e5	e, f	100101	e	e5, e7, e8	100111
e6	f, g	101110	f	e5, e6	001011
e7	e, g	100110	g	e6, e7	111000
e8	c, e	010100			

$$\text{XOR}_A = \text{XOR}(000011, 001010, 000101, 011000) = 010100 = \text{ID}(e8)$$



# The Linear MPC Model

---

- We need to extend the above idea for the case where the cut has multiple cross edges.
- If the cut has more than one edge, then all such edges appear only once in  $XOR_A$ .
- The resulting  $XOR_A$  may actually end up being the id of some edge that is not a cut edge!
- So, we need a couple more techniques.

# The Linear MPC Model

---

- Our first technique is to replace the id of vertices with  $O(\log n)$  bit random strings of  $\{0, 1\}$ .
- With this,  $\Pr(\text{There exists an edge } e, \text{XOR}_A = \text{id}(e)) \leq \sum_e \Pr(\text{XOR}_A = \text{id}(e)) \leq {}^nC_2 \times (1/2)^{12\log n} = 1/n^{10}$ .
- This low probability helps us tide over the trouble of using actual vertex identifiers.

# The Linear MPC Model

---

- For our second technique, we do the following.
- We do not know the number of edges crossing the cut.
- But, the number lies between 0 to  $n$ .
- So, we will try all possible estimates for the above number. All in parallel!
- The estimate is good if the actual number of edges crossing the cut is within  $\frac{1}{2}$  to 2 of the estimated number.
- In essence, we try estimates such as  $\{1, 2, 4, \dots, n\}$ .

# The Linear MPC Model

---

- Let  $k$  be the actual number of cut edges.
- Run each estimate  $k'$  in  $[1, 2, 4, \dots, n]$  in parallel.
- With a given  $k'$ , we mark edges with a probability of  $1/k'$ .
- Consider the  $k'$  such that  $k'/2 \leq k \leq k'$
- With a marking probability of  $1/k'$ , the expected number of cut edges marked is 1.
- We can reuse the case of identifying one cut edge if we get to mark just one of the  $k$  cut edges and no other edges are marked.
- Let us calculate the probability for the above event.

# The Linear MPC Model

---

- Let  $S$  be the set of cut edges marked.
- $\Pr(|S| = 1) = k \cdot (1/k') (1 - 1/k')^{(k-1)}$   
 $\geq (k'/2) \cdot (1/k') (1 - 1/k')^{(k-1)}.$   
 $\geq (1/2) \cdot (1/4)$  as  $1 - x \geq 4^{-x}.$   
 $\geq 1/8.$

# The Linear MPC Model

---

- A few more nitty-gritty details are needed beyond the two techniques.
- In particular, we will need multiple independent runs with each  $k'$ .
- With a success probability of at least  $1/8$ , Chernoff bounds tell us that  $O(\log n)$  runs will help us.
- The overall technique is called graph sketching.

# The Linear MPC Model

---

- Since the memory needed per sketch per node is  $O(\log n)$ , the linear MPC model can support the technique.
- Read the details from the posted resources.
- Putting together everything, we get:
- The MST of a graph  $G$  can be obtained in  $O(1)$  rounds in the linear memory MPC model.

# The Sub-linear MPC Model

---

- In the sub-linear MPC model, the only best known solution for the MST problem so far is to use the Boruvka algorithm.
- Recall that Boruvka's algorithm is what is used in the GHS algorithm also.
- The round complexity is  $O(\log n)$ .



# The Sub-Linear Memory Model

---

- We will now show how to improve the number of rounds from  $O(\log n)$  to  $O(\log D)$  where  $D$  is the diameter of the graph.
- There are two techniques that we will use.
- The first technique is called graph exponentiation.
- The second technique is graph labelling.

# The Sub-Linear MPC Model

---

- In this model, recall that each machine has a space of  $n^\varepsilon$  for some  $\varepsilon$ ,  $0 < \varepsilon < 1$ .
- Let us assume that each vertex has access to a memory of  $n^\varepsilon$ .
- Assume also that the graph is indeed a single connected component.
- In other words, we are verifying if the graph is connected.

# Graph Exponentiation

---

- Consider a graph which consists of  $n/D$  paths of length  $D$  where  $D \leq n^\epsilon$ .
- Within  $O(\log D)$  MPC rounds, every vertex can gather the entire neighborhood using graph exponentiation as follows.
- In every communication round, each node  $u$  informs its neighbors of the nodes contained in  $N(u)$ .

# Graph Exponentiation

---


- Consider a graph which consists of  $n/D$  paths of length  $D$  where  $D \leq n^\epsilon$ .
- Within  $O(\log D)$  MPC rounds, every vertex can gather the entire neighborhood using graph exponentiation as follows.
- In every communication round, each node  $u$  informs its neighbors of the nodes contained in  $N(u)$ .
- Then, every node can add the new nodes it learned about in its neighborhood by adding a virtual edge to each such node.

# Graph Exponentiation

---

- In every communication round, each node  $u$  informs its neighbors of the nodes contained in  $N(u)$ .
- Then, every node can add the new nodes it learned about in its neighborhood by adding a virtual edge to each such node.
- This way, in round  $j$  of the procedure, node  $u$  will be informed about all nodes and edges in its  $2^j$  - hop neighborhood.
- This procedure requires only  $O(\log D)$  rounds to obtain the  $D$ -hop neighborhood.

# Example : Graph Exponentiation



Node	Round 0	Round 1	Round 2	Round 3
a	b	a, c		
b	a, c	a, c, d,		
c	b, d	b, d, a, e	a, c, d, e, b, f, g	
d	c, e	c, e, b, f		
e	d, f	d, f, c, g		
f	e, g	d, e, g, h		
g	f, h	f, h, e		
h	g	f, h		

# Graph Exponentiation

---

- In every communication round, each node  $u$  informs its neighbors of the nodes contained in  $N(u)$ .
- Then, every node can add the new nodes it learned about in its neighborhood by adding a virtual edge to each such node. This way, in round  $j$  of the procedure, node  $u$  will be informed about all nodes and edges in its  $2^j$ -hop neighborhood.
- This procedure requires only  $O(\log D)$  rounds to obtain the  $D$ -hop neighborhood.
- The node can then compute connectivity of the collected subgraph in a single MPC round locally.
- Note that this will require a global memory of  $O(nD)$ .

# Graph Exponentiation

---

- In round  $j$  of the procedure, node  $u$  will be informed about all nodes and edges in its  $2^j$ -hop neighborhood.
- This procedure requires only  $O(\log D)$  rounds to obtain the  $D$ -hop neighborhood.
- The node can then compute connectivity of the collected subgraph in a single MPC round locally.
- Note that this will require a global memory of  $O(nD)$ .



# Graph Exponentiation

---



- If the  $D$ -hop neighborhood of a vertex fits into memory, then one can gather the entire neighborhood in  $O(\log D)$  MPC rounds.
- However, for general graphs, one cannot hope to collect the entire  $D$ -hop neighborhood of a vertex into a single machine.
  - This  $D$ -hop neighborhood can exceed the space in a machine.
- So, this technique works for graphs with low enough degree.

# Label Contraction

---

- The basic idea of label contraction is as follows.
- Mark each vertex independently with probability  $p = 1/2$ .
- For each unmarked vertex with a marked neighbor, relabel itself to one of the marked neighbor's label.
- By treating vertices with the same label as a supernode, we see that label contraction essentially contracts vertices in each round via relabelling.

# Label Contraction



# Label Contraction

---

- Claim: In expectation, at most  $\frac{3}{4}$  fraction of the existing labels remain after each iteration.
- Proof uses simple probability as follows.
- After the iteration, the only remaining labels are those of marked vertices.
- Consider a node  $v$  with a label  $l(v)$ .
- Each vertex is marked independently with probability  $p = 1/2$ .
- For an unmarked vertex  $v$ , the probability that it does not have a marked neighbor is  $(1/2)^{\deg(v)} \leq \frac{1}{2}$ .

# Label Contraction

---

- The probability that a vertex  $v$  is unmarked and has no marked neighbors is at most  $\frac{1}{4}$ .
  - Unmarked with a probability of  $\frac{1}{2}$  and has no marked neighbors with a probability of  $\frac{1}{2}$ .
- Hence,  $\Pr[ I(v) \text{ remains after one iteration} ]$   
     $= \Pr(v \text{ is marked, or } v \text{ is unmarked and has no marked neighbor})$   
     $\leq \Pr(v \text{ is marked}) + \Pr(v \text{ is unmarked and has no marked neighbor})$   
     $\leq \frac{1}{2} + \frac{1}{4} = \frac{3}{4}$ .
- The expectation follows.

# Label Contraction

---

- Can also use tail inequality to show that the above happens with high probability.
- Thus, in  $O(\log n)$  rounds, all end points of any edge share the same label.

# Combining the Two Techniques

---

- The exponentiation technique works best when the graph is sparse.
- The label propagation technique works best when the graph is dense.
  - Quick convergence of labels.
- Let us now see how to combine both of these techniques.

# Combining the Two Techniques

---

- The approach proceeds in phases.
- In each phase, the following happens.
  - From each node, perform graph exponentiation until the number of edges leaving the collected set of vertices exceeds  $n^{\epsilon/2}$ . Collapse the collected vertices into a supernode.
  - Perform label contraction by marking vertices independently with probability  $p = \log n / n^{\epsilon/2}$ .



# Combining the Two Technique

---

- The first step of each phase runs in  $O(\log D)$  rounds. (or fewer)
- The second step can be analyzed as earlier.
- Notice that after the first step, the degree of each (super)vertex is at least  $n^{\varepsilon/2}$ .
- So, the probability that a label does not get removed in a step is at most  $p + (1-p)(1-p)^{n^{\varepsilon/2}}$ .
- The above quantity is in  $O(p)$  with  $p = \log n/n^{\varepsilon/2}$ .
- We can deduce a high probability result so that in each step the number of labels drop by a factor of  $O(n^{\varepsilon/2})$ .
- The number of phases is therefore  $O(1/\varepsilon)$ .

# Sublinear MPC

---

- The above steps and techniques assumed the knowledge of the diameter  $D$  of the graph.
- This is not easy to obtain.
- One way to circumvent this difficulty is to try for multiple possible  $D$  values.
- Try with multiple estimates  $D'$  of  $D$  such that  $D'$  is equal to  $2^{2^i}$  for  $i = 1, 2, \dots$
- Run the algorithm for all  $D'$  in parallel.
- Two guidelines in finding what all possible values to try.
  - Safe: MPC constraints are not violated.
  - Checkable: Quickly check if the solution with some choice is correct.

# Summary of the MPC model and MST

---

- Super-linear memory:  $O(1/e)$  rounds, easy technique.
- Linear memory:  $O(1)$  rounds, but very complicated and lots of techniques.
- Sub-linear memory:  $O(\log D)$  rounds.

# MPC Model – Other Graph Algorithms

---

- Wide body of literature wrt MPC algorithms for various graph problems including MIS.
- Many of these results use randomization.
- Recent successes include derandomization as well.
- Variants of the MPC model also under study.
- A good thesis topic with lots of interesting problems yet to be solved.