



# MPI Tutorial

Aditya



# What is MPI?

- ❑ The Message Passing Interface (MPI) is a standard that defines an interface for message-passing libraries.
  - ❑ It is NOT a language, a compiler or a library.
- ❑ MPI is language-independent.
- ❑ OpenMPI is an open-source implementation of MPI in C, C++ and Fortran.
- ❑ mpi4py is a Python library that implements MPI.



# Why use MPI?

- ▢ Allows communication between processes, which is crucial for parallel computing.
  - ▢ Why is this crucial?
- ▢ Is powerful, portable, scalable, and widely used.

# Point-to-point Communication

```
int MPI_Send(const void* buf,          // what to send
             int count,                // number of elements to send
             MPI_Datatype datatype,    // type of elements
             int dest,                 // rank of destination
             int tag,                  // message tag
             MPI_Comm comm)            // communicator
;

int MPI_Recv(void* buf,                 // where to receive
             int count,                 // number of elements to receive
             MPI_Datatype datatype,    // type of elements
             int source,                // rank of source
             int tag,                  // message tag
             MPI_Comm comm,             // communicator
             MPI_Status* status)       // status object
;
```

# Group Communication – Send Identical Copies

```
int MPI_Bcast(void* buffer,           // what to broadcast
              int count,              // number of elements to broadcast
              MPI_Datatype datatype,  // type of elements
              int root,               // the rank of the process broadcasting the data
              MPI_Comm comm           // communicator
);
```

```
int MPI_Reduce(const void* sendbuf,   // what to send
               void* recvbuf,        // where to receive
               int count,             // number of elements to send/receive
               MPI_Datatype datatype, // type of elements
               MPI_Op op,             // operation to perform
               int root,              // the rank of the process receiving the result
               MPI_Comm comm          // communicator
);
```

# Group Communication – Partition & Transfer an Array

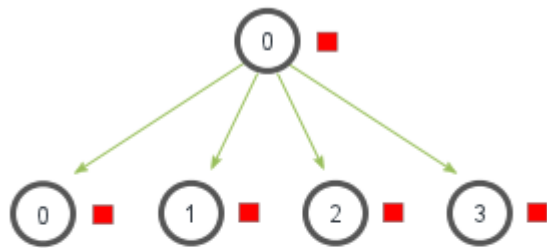
```
int MPI_Scatter(const void* sendbuf, // what to send
               int sendcount,       // number of elements to send to each process
               MPI_Datatype sendtype, // type of elements
               void* recvbuf,       // where to receive
               int recvcount,       // number of elements to receive
               MPI_Datatype recvtype, // type of elements
               int root,            // the rank of the process scattering the data
               MPI_Comm comm        // communicator
);

int MPI_Gather(const void* sendbuf, // what to send
               int sendcount,       // number of elements to send from each process
               MPI_Datatype sendtype, // type of elements
               void* recvbuf,       // where to receive
               int recvcount,       // number of elements to receive from each process
               MPI_Datatype recvtype, // type of elements
               int root,            // the rank of the process gathering the data
               MPI_Comm comm        // communicator
);
```

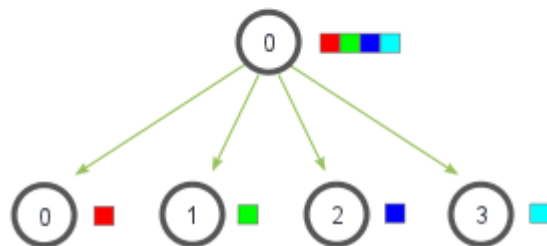
# Communication

- ❑ MPI\_Send, MPI\_Recv, MPI\_Bcast, MPI\_Reduce, MPI\_Scatter and MPI\_Gather are blocking calls.
- ❑ Senders and receivers must explicitly call these functions.

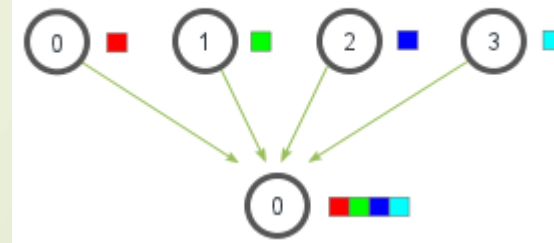
MPI\_Bcast



MPI\_Scatter



MPI\_Gather





# Hello, World!

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv); // initialises the MPI execution environment

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // gets the rank of this process
    MPI_Comm_size(MPI_COMM_WORLD, &size); // gets the number of processes

    printf("Hello, World! I am process %d out of %d processes.\n", rank, size);

    MPI_Finalize(); // terminates MPI execution environment
}
```



# Compile & Run

Compile:

```
mpicc hello-world.c -o hello-world
```

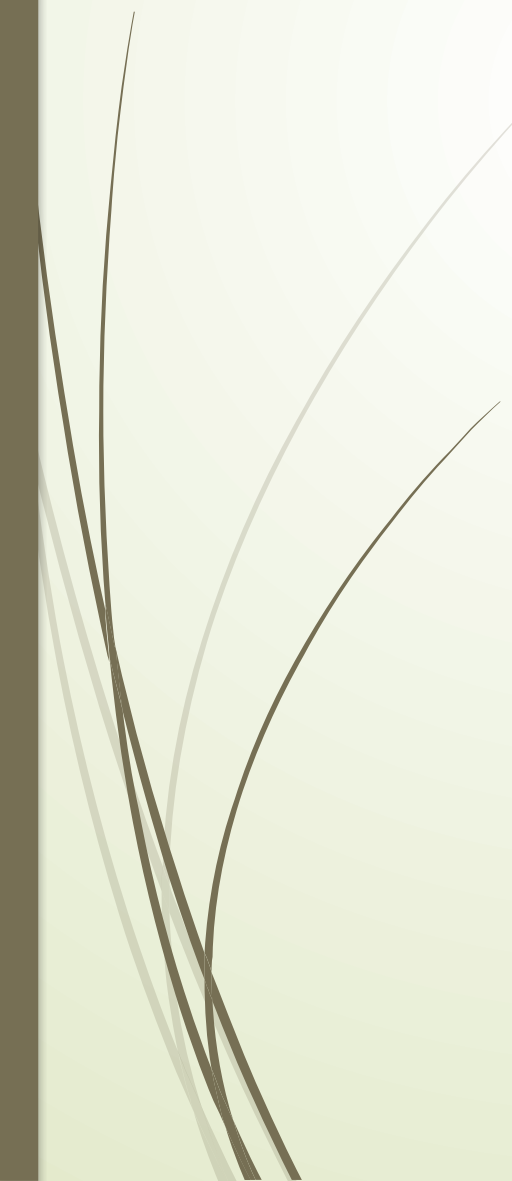
Run:

```
mpiexec -n 4 ./hello-world
```

```
• [aditya@indus mpi-tutorial]$ mpiexec -n 4 ./hello-world  
Hello, World! I am process 3 out of 4 processes.  
Hello, World! I am process 1 out of 4 processes.  
Hello, World! I am process 0 out of 4 processes.  
Hello, World! I am process 2 out of 4 processes.
```



# Suggested Resources

- ▮ <https://rookiehpc.github.io/mpi/docs/index.html>
  - ▮ <https://docs.open-mpi.org>
- 



Q&A